

1 Gradient Descent with Autograd

1.1 Implement gradient-based factorisation using PyTorch's AD


Due to length of code, we kindly ask the reader to visit the following [Github notebook](#) and find 'gd_factorise_ad' method for full implementation.

1.2 Factorise and compute reconstruction error on real data


With learning rate = 0.01 and 2000 epochs, reconstruction loss = 0.0254. Performing truncated SVD on the same data gives reconstruction loss = 1.0201.

1.3 Compare against PCA

SVD is a numerical method of computing the PCA. The reconstruction error is If mean-centred datapoint is \mathbf{X} and its approximations are $\tilde{\mathbf{X}}$ then reconstruction error is $\|\mathbf{X} - \tilde{\mathbf{X}}\|_F^2$. With Eckhart-Young Theorem, we know that an optimal choice of $\tilde{\mathbf{X}}$ is \mathbf{X}_k found via truncated SVD.



(a) From \mathbf{U} using truncated SVD



(b) From $\hat{\mathbf{U}}$ using factorisation

Figure 1: Projection of data onto principle directions.

2 A simple MLP

2.1 Implement the MLP

We kindly ask the reader to visit the following [Github notebook](#) and find 'mlp_train' method for full implementation.

2.2 Test the MLP




Figure 2: Left plot showing training loss. Right plot showing training accuracy in blue and validation accuracy in orange.